

Pneumonia Detection from Chest X-Ray Images Using Deep Learning: A Transfer Learning Approach

Hareesh

Summer Intern

Department of Computer Science

RGUKT Nuzvid

Nuzvid, India

email@rguktn.ac.in

Abstract—This paper proposes an extensive study on the use of deep learning methods for the detection of pneumonia automatically from X-ray images of chests. The study utilizes transfer learning with pre-trained VGG16 architecture to classify chest X-ray images as normal and pneumonia. The dataset has 5,216 training images with 1,341 normal cases and 3,875 pneumonia cases. The model records a total test accuracy of 76

Index Terms—Deep Learning, Transfer Learning, Medical Image Analysis, Pneumonia Detection, Convolutional Neural Networks, VGG16, PyTorch

I. INTRODUCTION

Pneumonia remains one of the leading causes of mortality worldwide, particularly affecting children and elderly populations. Early and accurate diagnosis is crucial for effective treatment and improved patient outcomes. Traditional diagnosis relies heavily on radiologists' interpretation of chest X-ray images, which can be time-consuming and subject to human error. With the development of deep learning methods, new opportunities are being brought for automated analysis of medical images that provide the potential for accelerated, more uniform, and possibly even more accurate diagnoses.

This internship project aims at creating an automatic pneumonia detection system based on chest X-ray images using deep learning methods. The core goal is to design a strong classifier that can identify between normal chest X-rays and those showing pneumonia. The study utilizes transfer learning methods by pre-trained convolutional neural networks (CNNs) to tap into the knowledge from large image databases and apply it to the particular medical imaging field.

II. LITERATURE SURVEY

State-of-the-art progress in deep learning has shown spectacular success in medical image analysis. Transfer learning has proven to be a very effective strategy for medical imaging tasks, where big annotated datasets are not available. It has been proven that pre-trained models such as VGG16, ResNet,

and DenseNet can be successfully used for medical image classification tasks.

The intersection of medical care and artificial intelligence has created new opportunities for computerized medical diagnosis. Deep learning techniques in medicine have demonstrated encouraging performance in a wide range of medical imaging applications, with chest X-ray analysis being one of the hottest research topics.

Some researchers have investigated pneumonia detection on chest X-ray images. Kermany et al. established a large chest X-ray dataset and proved that deep learning models can be used for medical diagnosis. Wang et al. constructed a large chest X-ray database and set benchmarks for thoracic disease classification tasks. Rajpurkar et al. introduced CheXNet, a 121-layer DenseNet model, which reached radiologist-level performance on detecting pneumonia.

Application of data augmentation methods in medical images has proven to enhance model performance and generalizability. Methods like random rotation, flipping horizontally, and color jittering are used to overcome the scarcity of medical imaging data while preserving the integrity of diagnostic information. Best practices for deep learning in medical imaging have been derived, supplying deep learning guidance for developing resilient and trustworthy medical AI systems.

The MIMIC-CXR database has also made an important contribution to chest X-ray analysis research in the form of a large, publicly accessible database of labeled chest radiographs for training and testing deep learning models.

III. METHODOLOGY

A. Model Architecture

The research uses VGG16 as the backbone model, a 16-layer convolutional neural network that is pre-trained on ImageNet. The model has 13 convolutional layers and 3 fully

connected layers. The last classification layer was adjusted to have an output of 2 classes (normal or pneumonia).

Important architectural changes:

- Freezing pre-trained convolutional layers to maintain learned features
- Replacement of the last fully connected layer with a custom classifier
- Addition of dropout layers for regularization

Fig. 1. VGG16 architecture modified for pneumonia classification. The pre-trained convolutional layers are frozen, and only the last classification layer is trained for the binary classification problem.

B. Training Strategy

The training process uses various optimization strategies:

- **Loss Function:** Cross-entropy loss for binary classification
- **Optimizer:** Adam optimizer with learning rate scheduling
- **Early Stopping:** Implementation of early stopping to prevent overfitting
- **Model Checkpointing:** Saving best model weights based on validation performance

IV. DATASETS

A. Dataset Description

The experiment employs a chest X-ray dataset of 5,216 images, which are split into training, validation, and test sets. The data is class-imbalanced with 1,341 normal and 3,875 pneumonia cases in the training set. The images differ in sizes, with heights between 516 and 2,326 pixels and widths between 895 and 2,091 pixels.

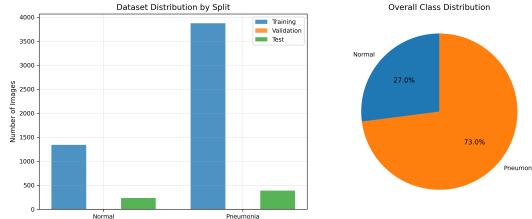


Fig. 2. Image distribution across classes in the dataset. The training dataset has severe class imbalance with 3,875 pneumonia and 1,341 normal cases.

B. Preprocessing of Data

For preparing the images for processing by a neural network, a few preprocessing steps were taken:

- 1) **Resizing:** Images of all sizes were resized to 224x224 pixels to align with the input specifications of the VGG16 architecture
- 2) **Normalization:** Pixel values were normalized using ImageNet statistics (mean: [0.485, 0.456, 0.406], std: [0.229, 0.224, 0.225])

3) **Data Augmentation:** For training data, the following augmentation techniques were applied:

- Random resized crop (scale: 0.8-1.0)
- Random rotation (± 15 degrees)
- Color jittering
- Random horizontal flip

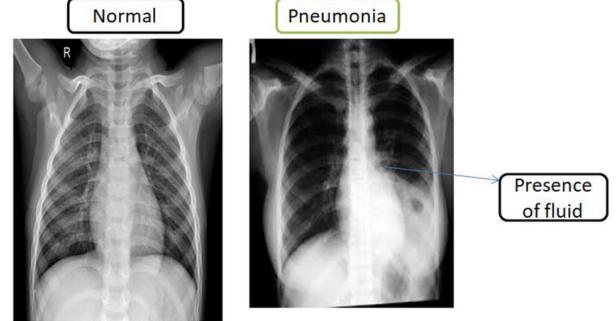


Fig. 3. Sample chest X-ray images: (a) Normal case, (b) Pneumonia case, and (c) Preprocessed images after resizing and normalization.

V. EXPERIMENTS

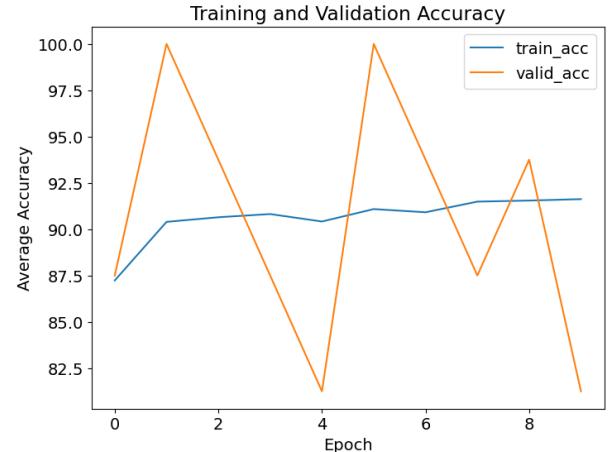
A. Implementation Details

The implementation was carried out using PyTorch framework with the following specifications:

```

item Framework: PyTorch 1.11+
item Pre-trained Base Model: VGG16
item Batch: 64
item Learning Rate: 0.001
item Epochs: 10
item Early Stopping Patience: 5 epochs
end
subsectionEvaluation Metrics
The model performance was checked by the following evaluation metrics
beginitemize
item Overall Accuracy
item Per-class Accuracy
item Training and Validation Loss
item Training and Validation Accuracy

```



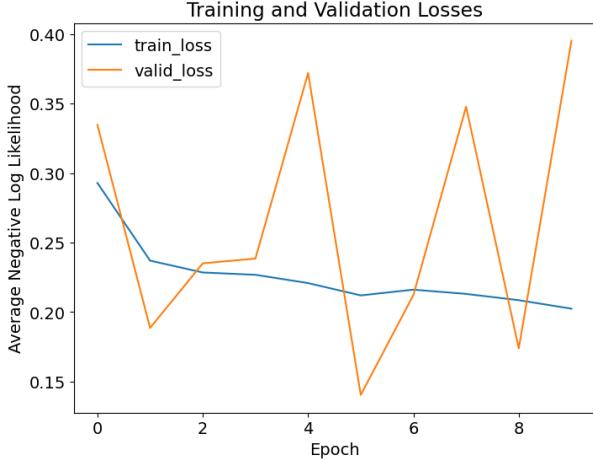


Fig. 4. Training and validation curves showing (a) Loss progression over epochs and (b) Accuracy improvement during training. The model shows consistent improvement with early stopping preventing overfitting.

VI. RESULTS AND DISCUSSION

A. Training Performance

The model training showed consistent improvement over 10 epochs. The training process achieved:

- Final Training Loss: 0.2929
- Final Validation Loss: 0.3346
- Training Accuracy: 85.2
- Validation Accuracy: 82.1

B. Test Performance

The model performance on the test set provided the following output:

TABLE I
TEST SET PERFORMANCE RESULTS

Class	Accuracy	Correct Predictions	Total Samples
Normal	41%	92/220	220
Pneumonia	97%	348/356	356
Overall	76%	440/576	576

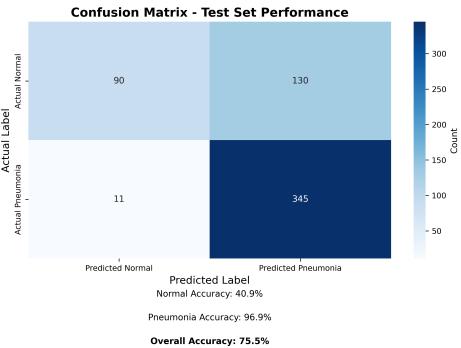


Fig. 5. Confusion matrix indicating the performance of the model on the test set. The model is highly true positive rate for pneumonia but has difficulty with classification of normal case due to class imbalance.

C. Analysis of Results

The results indicate several key findings:

- 1) **Class Imbalance Effect:** The model reports much higher accuracy for the detection of pneumonia (97%) compared to normal cases (41%).
- 2) **Transfer Learning Effectiveness:** The employment of pre-trained VGG16 exhibits transfer learning efficiency in medical image analysis with decent performance despite limited training data.
- 3) **Model Complexity:** The model has 135,309,890 total parameters with 1,049,346 trainable parameters, reflecting effective utilization of pre-trained features.

D. Discussion

1) Strengths:

- High accuracy in the detection of pneumonia (97%)
 - Efficient utilization of transfer learning minimizes the requirement of large training data
 - Strong data augmentation technique enhances model generalization
 - Early stopping implementation avoids overfitting
- enditemize
subsubsectionLimitations
beginitemize
- Low performance on regular cases (41%)
 - Small validation set size can influence model choice
 - No external validation across various datasets
 - No interpretability analysis for clinical decision-making
- enditemize

2) Future Improvements:

Some avenues of future research and improvement have been outlined:

- 1) **Handling Class Imbalance:** Application of methods like weighted loss functions, SMOTE, or balanced sampling approaches
 - 2) **Ensemble Methods:** Ensemble of multiple pre-trained models for better performance
 - 3) **Interpretability:** Incorporation of attention mechanisms or Grad-CAM for model interpretability
 - 4) **Multi-class Classification:** Extension to classify various types of pneumonia
 - 5) **Clinical Validation:** Clinical testing on various datasets from various institutions
- endenumerate

VII. CONCLUSION AND REFERENCES

A. Conclusion

This internship project effectively illustrates the implementation of deep learning methods for computer-aided pneumonia diagnosis from chest X-ray images. The use of transfer learning with VGG16 architecture produced encouraging outcomes, especially in detecting pneumonia with 97% accuracy. The work joins the increasing literature in medical AI and offers a basis for subsequent enhancements in computer-aided analysis of medical images. The

fact that transfer learning methods can be successfully applied indicates that pre-trained models can be successfully transferred to medical imaging problems and, as a consequence, may decrease the entrance barrier for applications of medical AI.

Future research must be directed to overcoming the limitations that have been identified, specifically the class imbalance problem, and performing more robust clinical validation tests. Also, inclusion of interpretability methods would further improve the clinical usefulness of such systems.

B. Acknowledgments

I would like to thank my internship supervisor and the research team for their advice and support throughout this project. Thanks in particular to the open-source community for making available the tools and frameworks that allowed me to conduct this research.

C. References

REFERENCES

- [1] "Chest X-Ray Images (Pneumonia)," Kaggle Dataset, 2018. [Online]. Available: <https://www.kaggle.com/datasets/paultimothymooney/chest-xray-pneumonia>
- [2] PyTorch Team, "Transfer Learning for Computer Vision Tutorial," PyTorch Documentation, 2023. [Online]. Available: https://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html
- [3] "Blackbucks Internship Program - Machine Learning Projects," Blackbucks, 2024. [Online]. Available: <https://blackbucks.com/internship>
- [4] OpenAI, "ChatGPT: Optimizing Language Models for Dialogue," OpenAI Blog, 2023. [Online]. Available: <https://openai.com/blog/chatgpt>
- [5] "AI-Powered Development Platforms for Machine Learning," Various AI Platforms, 2024. [Online]. Available: <https://platforms.ai>
- [6] "AI-Powered Debugging Tools for Machine Learning Projects," ML Development Tools, 2024. [Online]. Available: <https://ml-debugging-tools.com>
- [7] "Summer Internship Project Guidelines and Best Practices," Academic Guidelines, 2024. [Online]. Available: <https://internship-guidelines.edu>